

# *Fast Differentiable Rendering with 3D-GS*

*Based on slides from [Takikawa et al, 2023] and [Tulsiani, 2024]*

## **Outline**

- Differentiable Primitive Rendering
- Gaussian Splatting


## 3D Gaussian Splatting

## 3DGS: Differentiable Primitive Rendering




Slides from S. Tulsiani and V. Sitzmann

## Volumes: Rendering and Representation



Rendering Algorithm



(Tulsiani)

## Rendering Primitives (e.g. Gaussians)




1. Draw samples along the ray

2. Aggregate their contributions to render

(S. Tulsiani, 2024)


## From Ray tracing to Rasterization



Accelerating 3D Deep Learning with PyTorch3D. Ravi et al.

(S. Tulsiani, 2024)

## Differentiable Primitive Rendering



### Rasterization

```
def render(primitives, camera):
    """
    # Initialize a rasterization data structure
    # (records influencing primitives for each pixel)
    """

    for primitive in primitives:
        prim2d = project(primitive, camera)
        """
        # Update rasterization data structure
        """

    for pixel in camera.grid:
        """
        # Aggregate appearance from influencing primitives
        """
```

### Blending

(S. Tulsiani, 2024)

## Differentiable Gaussian Rendering

What is the representation  
of a 3D gaussian?

How to project to 2D  
and rasterize?

How to model/aggregate  
appearance?

Rasterization

{

```
def render(gaussians, camera):
    """
    # Initialize a rasterization data structure
    # (records influencing primitives for each pixel)
    """

    for gaussian in gaussians:
        gauss2d = project(gaussian, camera)
        """
        # Update rasterization data structure
        """

    for pixel in camera.grid:
        """
        # Aggregate appearance from influencing gaussians
        """
```

(S. Tulsiani, 2024)

Blending

{

## Differentiable Gaussian Rendering

What is the representation  
of a 3D gaussian?

How to project to 2D  
and rasterize?


How to model/aggregate  
appearance?



(S. Tulsiani, 2024)

## Differentiable Gaussian Rendering

**What is the representation of a 3D gaussian?**



**How to project to 2D and rasterize?**

$$\mathcal{G}_{\mathbf{V}}(\mathbf{x} - \mathbf{p}) = \frac{1}{2\pi|\mathbf{V}|^{\frac{1}{2}}} e^{-\frac{1}{2}(\mathbf{x}-\mathbf{p})^T \mathbf{V}^{-1} (\mathbf{x}-\mathbf{p})}$$

Factorize as scale and rotation:  $\mathbf{V} = RSS^T R^T$


Each gaussian also has an opacity and view-dependent color (via SH coefficients):  $\alpha, \mathbf{c}$

**How to model/aggregate appearance?**

(S. Tulsiani, 2024)


## Differentiable Gaussian Rendering

**What is the representation of a 3D gaussian?**



(S. Tulsiani, 2024)

**How to project to 2D and rasterize?**



$\mathbf{p}', R', S$

**How to model/aggregate appearance?**

We can use the camera extrinsics to transform each 3D gaussian to the camera frame

(S. Tulsiani, 2024)


## Differentiable Gaussian Rendering

What is the representation of a 3D gaussian?



$\mathbf{p}', R', S$

How to project to 2D and rasterize?



Q: What is the image-space projection of a 3D gaussian?

A: Can approximate as a 2D gaussian!

(EWA Volume Splatting, Zwicker et al., 2001)

How to model/aggregate appearance?

$$\pi(\mathbf{x}) = \mathbf{u} \quad z \begin{bmatrix} u \\ v \\ 1 \end{bmatrix} = K \begin{bmatrix} X \\ Y \\ Z \\ 1 \end{bmatrix}$$

$\pi$ : Projection function for mapping 3D points to pixels

2D mean:  $\mu_{2D} = \pi(\mu_{3D})$

2D covariance:

$$J = \frac{\partial \pi}{\partial \mathbf{x}}(\mu_{3D})$$

$$\Sigma_{2D} = J \Sigma_{3D} J^T$$


(S. Tulsiani, 2024)

## Differentiable Gaussian Rendering

What is the representation of a 3D gaussian?



How to project to 2D and rasterize?



$\mu_{2D} = \pi(\mu_{3D})$

$$\Sigma_{2D} = J \Sigma_{3D} J^T$$

How to model/aggregate appearance?

1. Sort gaussians from closest to furthest from the camera

2. For each pixel  $\mathbf{u}$ , compute opacity for each gaussian  $\mathcal{G}_k$ :

$$\bar{\alpha}_k = \alpha_k \frac{e^{-(\mathbf{u}-\mu_{2D}^k)^T (\Sigma_{2D}^k)^{-1} (\mathbf{u}-\mu_{2D}^k)}}{2\pi|\Sigma_{2D}^k|^{0.5}}$$

(In practice, can rasterize 'blocks' instead of entire image as not all gaussians influence all blocks)

(S. Tulsiani, 2024)

## Differentiable Gaussian Rendering

## What is the representation of a 3D gaussian?

# How to project to 2D and rasterize?

## How to model/aggregate appearance?



Compute per-gaussian weights based on opacities of current and previous gaussians:

$$w_k = \bar{\alpha}_k \prod_{j=1}^{k-1} (1 - \bar{\alpha}_j)$$

Use per-gaussian SH coefficients and ray direction to get view-dependent color  $\mathbf{c}_k$


Aggregate to obtain pixel color:

$$\mathbf{c} = \sum_k w_k \mathbf{c}_k$$

(S. Tulsiani, 2024)

## Computational Graph (gsplat)

- Forward ( $\uparrow$ ) and Backward ( $\downarrow$ ) Gaussian Splatting Rendering Function



## Properties of Gaussians for Rendering

Gaussians are closed under affine transforms, integration

$$\mathcal{G}_V(\mathbf{x} - \mathbf{p}) = \frac{1}{2\pi|V|^{\frac{1}{2}}} e^{-\frac{1}{2}(\mathbf{x}-\mathbf{p})^T V^{-1}(\mathbf{x}-\mathbf{p})}$$

↓  
3D Covariance!

Affine mapping  $\Phi = Mx + p$  of coordinates (such as cam2world matrix):

$$\mathcal{G}_V(\Phi^{-1}(\mathbf{u}) - \mathbf{p}) = \frac{1}{|M^{-1}|} \mathcal{G}_{MV M^T}(\mathbf{u} - \Phi(\mathbf{p}))$$

Integrate along axis:

$$\int_{IR} \mathcal{G}_V^3(\mathbf{x} - \mathbf{p}) dx_2 = \mathcal{G}_{\hat{V}}^2(\hat{\mathbf{x}} - \hat{\mathbf{p}})$$

$$V = \begin{pmatrix} a & b & c \\ b & d & e \\ c & e & f \end{pmatrix} \Leftrightarrow \begin{pmatrix} a & b \\ b & d \end{pmatrix} = \hat{V}$$

(V. Sitzmann, 2024)



•



•

•

•

## Transform Gaussians into Camera Coordinates

**Cam2world** is affine mapping  $\phi(x) = Wx + p$ :

$$\mathcal{G}_{V'_k}(\varphi^{-1}(\mathbf{u}) - \mathbf{t}_k) = \frac{1}{|W^{-1}|} \mathcal{G}_{V_k}(\mathbf{u} - \mathbf{u}_k) = r'_k(\mathbf{u})$$

**Projection**  $m(u)$  is not an affine mapping ∵


$$\begin{pmatrix} x_0 \\ x_1 \\ x_2 \end{pmatrix} = m(\mathbf{u}) = \begin{pmatrix} u_0/u_2 \\ u_1/u_2 \\ \| (u_0, u_1, u_2)^T \| \end{pmatrix}$$

$$\begin{pmatrix} u_0 \\ u_1 \\ u_2 \end{pmatrix} = m^{-1}(\mathbf{x}) = \begin{pmatrix} x_0/l \cdot x_2 \\ x_1/l \cdot x_2 \\ 1/l \cdot x_2 \end{pmatrix},$$


But can approximate with first-order Taylor Expansion as:

$$m_{u_k}(\mathbf{u}) = \mathbf{x}_k + J_{u_k} \cdot (\mathbf{u} - \mathbf{u}_k) \quad J_{u_k} = \frac{\partial m}{\partial u}(\mathbf{u}_k)$$

(V. Sitzmann, 2024)



## Transform Gaussians into Camera Coordinates



But can approximate with first-order Taylor Expansion as:

$$\mathbf{m}_{\mathbf{u}_k}(\mathbf{u}) = \mathbf{x}_k + \mathbf{J}_{\mathbf{u}_k} \cdot (\mathbf{u} - \mathbf{u}_k) \quad \mathbf{J}_{\mathbf{u}_k} = \frac{\partial \mathbf{m}}{\partial \mathbf{u}}(\mathbf{u}_k)$$

Projected, 2D Gaussians are then:

$$\frac{1}{|\mathbf{W}^{-1}| |\mathbf{J}^{-1}|} \mathcal{G}_{\mathbf{V}_k}(\mathbf{x} - \mathbf{x}_k)$$

$$\begin{aligned} \mathbf{V}_k &= \mathbf{J} \mathbf{V}'_k \mathbf{J}^T \\ &= \mathbf{J} \mathbf{W} \mathbf{V}''_k \mathbf{W}^T \mathbf{J}^T. \end{aligned}$$


Finally, can integrate along rays:

$$\begin{aligned} q_k(\hat{\mathbf{x}}) &= \int_{\mathbb{R}} \frac{1}{|\mathbf{J}^{-1}| |\mathbf{W}^{-1}|} \mathcal{G}_{\mathbf{V}_k}(\hat{\mathbf{x}} - \hat{\mathbf{x}}_k, x_2 - x_{k2}) dx_2 \\ &= \frac{1}{|\mathbf{J}^{-1}| |\mathbf{W}^{-1}|} \mathcal{G}_{\mathbf{V}_k}(\hat{\mathbf{x}} - \hat{\mathbf{x}}_k) \end{aligned}$$

(V. Sitzmann, 2024)

## Advantage of Rasterization

***Can compute volume rendering integral without ever sampling a single 3D point in space!***




$$C = \sum_{i \in \mathcal{N}} c_i \alpha_i \prod_{j=1}^{i-1} (1 - \alpha_j), \quad (3)$$

where  $c_i$  is the color of each point and  $\alpha_i$  is given by evaluating a 2D Gaussian with covariance  $\Sigma$  [Yifan et al. 2019] multiplied with a learned per-point opacity.


(V. Sitzmann, 2024)

## Problem: Local minima...



(V. Sitzmann, 2024)

## Gaussian Splatting: Bells and Whistles




Fix 1:  
Initialize with sparse point cloud from SfM

Fix 2:  
Split/clone gaussians  
based on heuristics

3D Gaussian Splatting for Real-Time Radiance Field Rendering. Kerl et. al.


(S. Tulsiani, 2024)

## Fix 1: Start from SFM point cloud. (Initialization)



(V. Sitzmann, 2024)

## Fix 2: Heuristic pruning and spawning operations (Adaptation)



(V. Sitzmann, 2024)

**“... And Many More Details !”**

*- LV*