


Volume Differentiable Rendering & NeRFs

Based on slides from [Gkioulekas, 2025], [Takikawa et al, 2023] and [Tulsiani, 2024]

Recap - Volume Rendering Pipeline

- Differentiable Volumetric Rendering Function



- Neural Volumetric 3D Scene Model

$$(x, y, z, \theta, \phi) \rightarrow F_{\Omega} \rightarrow (r, g, b, \sigma)$$

- Reconstruction via Analysis-bySynthesis



Outline


- Radiative Transfer Equation
- Volume Rendering Equation
- Differentiable Rendering
- 3D Scene Models
- NeRFs

Radiative Transfer Equation

Slides from [Gkioulekas, 2025]


Participating Media

Typically, we do not model particles of a medium explicitly
(wouldn't fit in memory, completely impractical to ray trace)




The properties are described statistically using various coefficients and densities

- Conceptually similar idea as microfacet models




Heterogeneous (spatially varying coefficients):

- Procedurally, e.g., using a noise function
- Simulation + volume discretization, e.g., a voxel grid




Differential Beam Segment




(Gkioulekas)

Radiative Transfer Equation (RTE)



(Gkioulekas)

Emission



$$dL(\mathbf{x}, \vec{\omega}) = \sigma_a(\mathbf{x})L_e(\mathbf{x}, \vec{\omega})dz$$


*Sometimes modeled without the absorption coefficient just by specifying a “source” term

$\sigma_a(\mathbf{x})$: absorption coefficient $[m^{-1}]$

$L_e(\mathbf{x}, \vec{\omega})$: emitted radiance

(Gkioulekas)

Absorption




$$dL(\mathbf{x}, \vec{\omega}) = -\sigma_a(\mathbf{x})L(\mathbf{x}, \vec{\omega})dz$$

$\sigma_a(\mathbf{x})$: absorption coefficient $[m^{-1}]$

(Gkioulekas)

Out-scattering




$$dL(\mathbf{x}, \vec{\omega}) = -\sigma_s(\mathbf{x})L(\mathbf{x}, \vec{\omega})dz$$

$\sigma_s(\mathbf{x})$: scattering coefficient $[m^{-1}]$

(Gkioulekas)

In-scattering



$$dL(\mathbf{x}, \vec{\omega}) = \sigma_s(\mathbf{x}) L_s(\mathbf{x}, \vec{\omega}) dz$$

$\sigma_s(\mathbf{x})$: scattering coefficient $[m^{-1}]$

$L_s(\mathbf{x}, \vec{\omega})$: in-scattered radiance

(Gkioulekas)

Complexity Progression - Scattering

homogeneous vs. heterogeneous

- none
- fake ambient
- single
- multiple

(Gkioulekas)

Volume Rendering Equation

Slides from [Gkioulekas, 2025]

Transmittance


Homogeneous volume:

$$T_r(\mathbf{x}, \mathbf{y}) = e^{-\sigma_t \|\mathbf{x} - \mathbf{y}\|}$$

Heterogeneous volume (spatially varying σ_t):

$$T_r(\mathbf{x}, \mathbf{y}) = e^{-\int_0^{\|\mathbf{x} - \mathbf{y}\|} \sigma_t(t) dt}$$

Optical thickness



Volume Rendering Equation


$$\begin{aligned} L(\mathbf{x}, \vec{\omega}) = & T_r(\mathbf{x}, \mathbf{x}_z) L(\mathbf{x}_z, \vec{\omega}) \\ & + \int_0^z T_r(\mathbf{x}, \mathbf{x}_t) \sigma_a(\mathbf{x}_t) L_e(\mathbf{x}_t, \vec{\omega}) dt \\ & + \int_0^z T_r(\mathbf{x}, \mathbf{x}_t) \sigma_s(\mathbf{x}_t) \int_{S^2} f_p(\mathbf{x}_t, \vec{\omega}', \vec{\omega}) L_i(\mathbf{x}_t, \vec{\omega}') d\vec{\omega}' dt \end{aligned}$$

(Gkioulekas)

Volume Rendering Equation

$$L(\mathbf{x}, \vec{\omega}) = \boxed{T_r(\mathbf{x}, \mathbf{x}_z) L(\mathbf{x}_z, \vec{\omega})}$$

\downarrow
Reduced (background) surface radiance




(Gkioulekas)

Volume Rendering Equation

$$L(\mathbf{x}, \vec{\omega}) = T_r(\mathbf{x}, \mathbf{x}_z)L(\mathbf{x}_z, \vec{\omega}) + \int_0^z T_r(\mathbf{x}, \mathbf{x}_t)\sigma_a(\mathbf{x}_t)L_e(\mathbf{x}_t, \vec{\omega})dt$$

Accumulated emitted radiance




(Gkioulekas)

Volume Rendering Equation

$$L(\mathbf{x}, \vec{\omega}) = T_r(\mathbf{x}, \mathbf{x}_z)L(\mathbf{x}_z, \vec{\omega}) + \int_0^z T_r(\mathbf{x}, \mathbf{x}_t)\sigma_a(\mathbf{x}_t)L_e(\mathbf{x}_t, \vec{\omega})dt + \int_0^z T_r(\mathbf{x}, \mathbf{x}_t)\sigma_s(\mathbf{x}_t) \int_{S^2} f_p(\mathbf{x}_t, \vec{\omega}', \vec{\omega})L_i(\mathbf{x}_t, \vec{\omega}')d\vec{\omega}'dt$$

Accumulated in-scattered radiance




(Gkioulekas)

Differentiable Rendering

Slides from [Tulsiani, 2024]

Differentiable Volume Rendering

- Volume Rendering Setting




$$\propto L(x^*, \omega)$$

radiance for: x^* = pixel sensor centre, w = direction from x to optical centre

(Tulsiani)


Mathematical Model

- Emission-Absorption Volume Rendering



Mathematical Model (simplified)

- Emitted Radiance from Inside Volume




$$L(\mathbf{x}, \omega) = \int_0^z T(\mathbf{x}, \mathbf{x}_t) \sigma(\mathbf{x}_t) L_e(\mathbf{x}_t, \omega) dt$$

Only the Accumulated Emitted Radiance Term!

(Tulsiani)

Computational Volume Rendering: Ray Marching



$$L(\mathbf{x}, \omega) = \sum_{i=1}^N T(\mathbf{x}, \mathbf{x}_{t_i}) (1 - e^{-\sigma_{t_i} \Delta t}) L_e(\mathbf{x}_{t_i}, \omega)$$


Multiplicativity

$$T(\mathbf{x}, \mathbf{x}_{t_i}) = T(\mathbf{x}, \mathbf{x}_{t_{i-1}}) e^{-\sigma_{t_{i-1}} \Delta t}$$

Assume constant coefficient
between samples

(Tulsiani)

Computational Volume Rendering: Ray Marching



1. Draw uniform samples along a ray (N segments, or N+1 points)
2. Compute transmittance between camera and each sample
3. Aggregate contributions across segments to get overall radiance (color)

(Tulsiani)

Rendering Model

- Computation for a Ray

Rendering model for ray $\mathbf{r}(t) = \mathbf{o} + t\mathbf{d}$:

$$\mathbf{c} \approx \sum_{i=1}^n T_i \alpha_i \mathbf{c}_i$$

differentiable w.r.t. \mathbf{c}, σ


colors
weights

How much light is blocked earlier along ray:

$$T_i = \prod_{j=1}^{i-1} (1 - \alpha_j)$$

How much light is contributed by ray segment i :

$$\alpha_i = 1 - \exp(-\sigma_i \delta_i)$$



(Tulsiani)

Computational Volume Rendering: A summary

$$\sigma_{t_i} \equiv \sigma(\mathbf{x}_{t_i}) \longrightarrow L(\mathbf{x}, \omega)$$

$$L_e(\mathbf{x}_{t_i}, \omega)$$

If we can compute:

- (per-point) density
 - (per-point, direction) emitted light,
- we can render **any** ray through the medium


Equivalently, we can render an image from any camera viewpoint (using H^*W rays)

Note: **Differentiable** process w.r.t. the **density, emitted light**

and also camera parameters if density, emission are differentiable functions of position, direction

(Tulsiani)

Volumes: Rendering and Representation



Rendering Algorithm



How to represent volumes?

(such that we can compute pointwise density and emitted light)



(Tulsiani)

Modeling the Scene

Slides from [Takikawa et al, 2023]]

3D Scene Model


- Density Fields




$$V[x, y, z] \in \mathbf{R}^+ \quad f(\mathbf{p}) \in \mathbf{R}^+$$

(Tulsiani)

Volumes: Rendering and Representation




Rendering Algorithm



Option 1: A grid

(Tulsiani)


Volumes: Rendering and Representation




(Tulsiani)

Neural 3D Scene Models

- Signed Distance Function



- Radiance Field



(Takikawa)

NeRF


Slides from[Takikawa et al, 2023]]

Neural Radiance Fields

- Characteristics

Strengths:


Compactness



Self-regularizing

$$\operatorname{argmin}_x \|y - F(x)\| + \lambda P(x).$$

Domain agnostic



Weaknesses:

Computationally expensive

Not easily editable


Hard to model semantics and discrete data

Lack of theoretical understanding

(Takikawa)


Analysis-by-Synthesis Reconstruction

- General Framework




Reconstruction Method

- Elements



- Algorithm




Loss Function

- Optimization


NeRF Innovation

NeRF uses volume rendering
(i.e. removes the Dirac assumption
in Kajiya's rendering equation)



Forward Map (e.g. differentiable rendering)

$$\min_{\theta} \parallel F(\begin{matrix} x \\ y \\ z \end{matrix}, d) - f_{\theta}(x, y, z) = d \parallel$$




Baking via solving inverse problems!

(Takikawa)

Learning Neural Radiance Fields



- 1) Acquire multiple images of a scene with associated camera viewpoints



2) Design a neural network $(\sigma, c) = f_{\theta}(x, \omega)$


$$\min_{\theta} \sum_I \sum_{\mathbf{p}} \| \text{render}(\mathbf{p}, \pi; f_{\theta}) - I[\mathbf{p}] \|^2$$

- 3) Train with a view-synthesis loss using volume rendering

Learning Neural Radiance Fields

$$\min_{\theta} \sum_I \sum_{\mathbf{p}} \| \text{render}(\mathbf{p}, \pi; f_{\theta}) - I[\mathbf{p}] \|_2^2$$

(pixel, camera) -> ray
volume rendering



NeRF: Representing Scenes as Neural Radiance Fields for View Synthesis.

(Tulsiani)

Learning Neural Radiance Fields



100 training images

Details: Spectral Projection
(Fourier Feature Mapping)



(using position encodings in input)



Novel-view Renderings

A great example that 'execution matters'

NeRF: Representing Scenes as Neural Radiance Fields for View Synthesis.

(Tulsiani)

Foreground and Background Radiance

NeRF++: ANALYZING AND IMPROVING NEURAL RADIANCE FIELDS


Kai Zhang
Cornell Tech

Gernot Riegler
Intel Labs



Noah Snavely
Cornell Tech

Vladlen Koltun
Intel Labs

$$B = \{(x, y, z) : \sqrt{x^2 + y^2 + z^2} = 1\}$$



Unbounded Scenes




A ‘background’ NeRF uses normalized sphere coordinate and $1/r$ as input

Uniform sampling for ray segments inside sphere, and $1/r$ based sampling outside

(Tulsiani)


View-dependent Effects




Same 3D point — different color based on viewing direction

Do we need Neural Radiance?

From NeRF to Grids & S.H.



Spherical Harmonics: A basis for scalar functions on a sphere



View-dependent color can be inferred via basis coefficients

Plenoxels: Radiance Fields without Neural Networks. Yu et. al.

(Tulsiani)